

Django-MySQL

Adam Johnson - me@adamj.eu

12 May 2015

Database Share

- ▶ Who's using PostgreSQL?
- ▶ Who's using MySQL?

Database Share

- ▶ Worldwide, of open source databases:
 - ▶ MySQL 56%
 - ▶ MariaDB 18% (MySQL fork!)
 - ▶ PostgreSQL 13%
- ▶ (scalebase.com)

Motivation

- ▶ Django 1.8 comes with “django.contrib.postgres”
- ▶ Hey, MySQL should have something too!
- ▶ Separate library = more thorough testing, multiple django versions supported

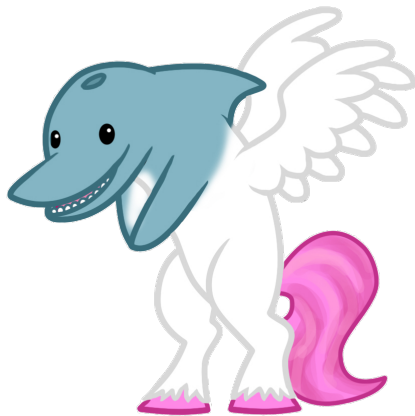
Mascot



+



Mascot



It can't be unseen.

1. Database Functions

- ▶ New in Django 1.8!
- ▶ "The" place to find MySQL-only functions

```
>>> Author.objects.annotate(  
...     full_name=ConcatWS('first_name',  
...                        'last_name',  
...                        separator=' ')  
... )  
>>> Author.objects.order_by(  
...     Field('gender', ['male', 'female'])  
... )
```

2. Named Locks

- ▶ Easy way to limit access to a resource
- ▶ E.g. API with connection limit

```
try:  
    with Lock('my_unique_name', acquire_timeout=2.0):  
        mutually_exclusive_process()  
except TimeoutError:  
    print "Could not get the lock"
```


3. Approximate Count

- ▶ `Model.objects.count()` = `SELECT COUNT(*)` which requires a table scan - slow!
- ▶ Various snippets and libraries exist, but this is (hopefully) *the bestTM*.
- ▶ Easy to hook into external code such as admin

```
>>> Author.objects.count() # slow
509741
>>> Author.objects.approx_count() # fast , some error
531140
```

4. List and Set Fields

- ▶ Cousins to `django.contrib.postgres's ArrayField`
- ▶ Store values in comma-separated strings, lookups and F implementations using MySQL functions.

```
# models
class Person(Model):
    post_nominals = ListTextField(
        base_field=CharField(max_length=32)
    )

# shell
>>> Person.objects.create(
...     name='Horatio',
...     post_nominals=['PhD', 'Esq.', 'III']
... )
>>> Person.objects.filter(post_nominals__contains='PhD')
[<Person: Horatio>]
```

5. Smart Iteration

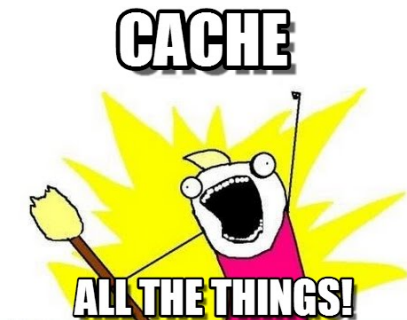
- ▶ Iterate and modify "big data" without fear, in primary-key-based slices
- ▶ Dynamically modifies slice size and checks MySQL status to avoid long-running outage-causing operations

```
# Turn this ...
min_id, max_id = 0, 1000
max_author_id = Author.objects.order_by('-id')[0].id
while True:
    author_slice = Author.objects.filter(
        address="Nowhere",
        id__gte=min_id,
        id__lte= # BLA BLA BLA
    # YOU GET THE IDEA IT'S A LOT OF CODE

# ...into this:
bad_authors = Author.objects.filter(address="Nowhere")
for author in bad_authors.iter_smart():
    author.send_apology_email()
```

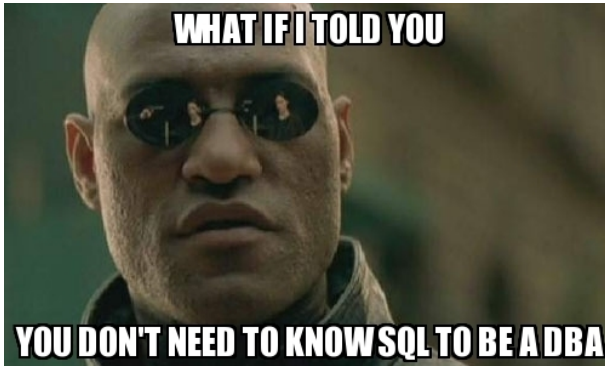
6. Cache Backend

- ▶ DatabaseCache is slow, too generic
- ▶ Make use of MySQL's upsert syntax so that each operation only takes one query
- ▶ Does automatic compression with zlib, like pylibmc
- ▶ And many other improvements



7. Migration Operations

- ▶ **InstallPlugin** - Installs a plugin
- ▶ **InstallSOName** - Installs every plugin in a library file
- ▶ **AlterStorageEngine** - Alters a table's storage engine



Coming Soon...

- ▶ Dynamic Columns - MariaDB's answer to HStore/json but with extra types
- ▶ Query hints in the ORM like **STRAIGHT_JOIN**
- ▶ Migration -¿ **pt-online-schema-change** rewriter
- ▶ 37 Github Issues open!





- ▶ github.com/adamchainz/django-mysql
- ▶ Contributors wanted!
- ▶ me@adamj.eu